SPRINT PLAN · FOUR PEOPLE · THREE DAYS

Profile Intelligence Working Slice

One niche, three sources, end to end, running on real people

A freelancer uploads a CV, gives a GitHub username and pastes their existing profile text. They get back a readiness score, a ranked list of gaps, and a rewritten title and overview where every number can be traced to where it came from. That is the whole target.

Task division, handoff contracts, checkpoints and the scope fence

1. What gets built

Not the full system. A thin vertical slice through it, working end to end, on one niche.

The demo path. If this runs on five real people, the sprint succeeded.

Everything in the specification that is not on this path is deliberately out. The point of three days is to find out whether the gap report tells a working freelancer something useful, and that question can be answered without automated benchmarks, without OCR, and without a login screen.

2. Who owns what

Owner	Area	Ships
M1	Ingestion	File router, CV parsing, GitHub pull, Upwork text paste. Everything that turns raw input into claim records.
M2	Benchmark and evidence	The hand-built benchmark for one niche, the required-terms list, and the rules that assign a tier to each claim.
M3	Scoring and generation	Seven dimensions, both caps, gap ranking with the blocking and dependency rules, title and overview generation, span validation.

Owner	Area	Ships
M4	Platform and integration	Storage, API, the page that displays results, and wiring the other three together as their parts land.

M4 has a second job that matters more than the first: being the integrator. On a three-day build the failure mode is never that someone did not finish their piece, it is that three finished pieces do not fit together on the last afternoon.

3. The schedule

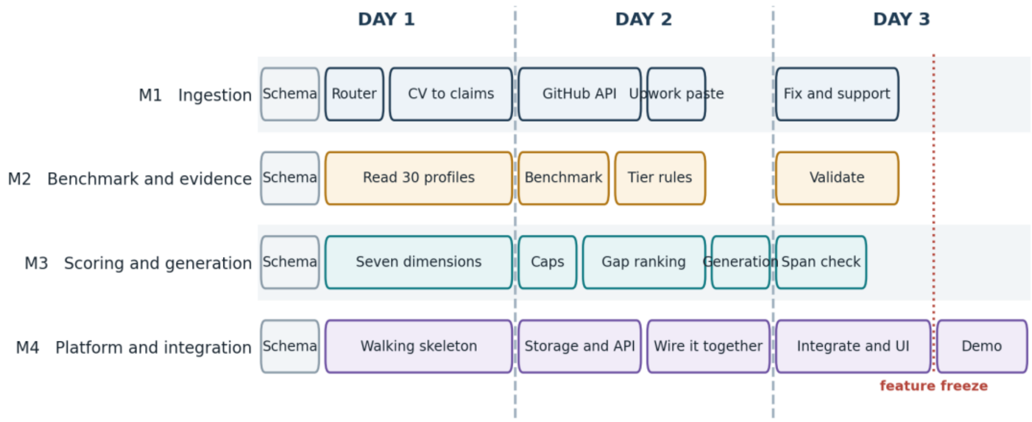

Four lanes, three days. The grey block at the start is the only thing everyone does together.

The critical path is two hours long

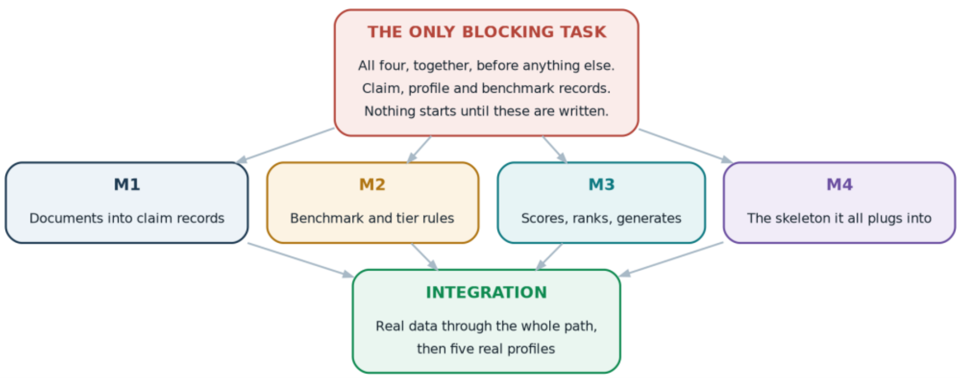

One blocking task. After it, four people work in parallel and nobody waits.

Spend the first two hours writing down three record shapes

The claim record: what a single piece of evidence looks like, with its source span, tier, weight and date.

The profile record: how claims aggregate into skills, and what the scorer reads.

The benchmark record: required terms, per-dimension targets and rate bands for the chosen niche.

Write them as actual files, not as a conversation. Everyone codes against them from that moment, using fake data until the real thing arrives.

After midday on day one, changing a schema needs all four people to agree. This sounds heavy and it is the single rule most likely to save the sprint. Schema churn on day two is what turns a three-day build into a five-day build.

4. The scope fence

IN SCOPE

- One niche only
- CV upload, GitHub, Upwork paste
- Hand-built benchmark, thirty profiles
- Eight evidence tiers, rules based
- Seven dimensions with both caps
- Gap ranking with blocking rules
- Title and overview generation
- One page that shows the result

OUT OF SCOPE

- Accounts and login, hardcode one user
- OCR, native PDFs only, reject scans loudly
- LinkedIn, portfolio, video, articles
- Automated benchmark scraping
- Skill taxonomy mapping
- Deduplication and entity resolution
- Weight refit, presence engine
- Proposal drafts, multiple niches
- Any visual polish at all

The right column is what makes the left column possible in three days

The out-of-scope list is not a wish list for later. It is a commitment for now. If someone finishes early, they help with integration or run more real profiles through the system. They do not start on OCR.

Two specific traps

CV parsing is a rabbit hole. Support a handful of common layouts and fail loudly on anything else, with a message telling the user what went wrong. Do not spend day two chasing a two-column PDF with a photo in it.

The benchmark is manual reading, and manual reading always takes longer than people estimate. Cap it at thirty profiles in one niche. Thirty is enough to see the pattern; the statistics can improve later.

5. Day by day

Day one

Goal: a skeleton that runs the whole path on fake data, and everyone unblocked.

Owner	Task	Done when
All	Lock the claim, profile and benchmark record shapes. Write them as files.	Three schema files committed, all four agree
M1	File type router, then CV parsing into claim records	A real CV produces real claim records
M2	Read thirty top profiles in the chosen niche and pull out the patterns by hand	Notes on title shapes, overview length, recurring terms, rate bands
M3	The seven dimension formulas, running against fake claims	Fake input produces a plausible score breakdown
M4	Walking skeleton: API, page, and stub endpoints returning fake data	The whole path runs from a browser, with nothing real in it

The skeleton matters more than it looks. By the end of day one there should be a URL that takes an input and shows a score, even if every number in it is invented. From then on, every real piece replaces a fake one, and integration stops being a cliff at the end.

Day two

Goal: one real person, all the way through.

Owner	Task	Done when
M1	GitHub API pull, then Upwork profile text paste parsing	Three sources all produce claim records in the agreed shape
M2	Benchmark written as a file, required terms listed, tier rules implemented	Any claim can be assigned a tier automatically
M3	Caps applied, gap ranking with blocking and dependency rules, then generation	A score object plus a ranked gap list plus a title and overview
M4	Real storage, real API, replace stubs with the pieces as they land	Fake data is gone from the main path

End of day two is the real checkpoint. One team member's own CV and GitHub should go in and produce a score, a ranked gap list and a rewritten title. If that is not working by the end of day two, day three becomes triage rather than polish, and the demo scope shrinks.

Day three

Goal: five real profiles, then stop.

Owner	Task	Done when
M1	Fix whatever breaks on real inputs, support the others	Five different CVs parse without crashing
M2	Check the benchmark against real profiles, adjust required terms	Scores look sane to someone who knows the niche
M3	Span validation on every generated claim, then fixes	No generated number exists without a traceable source
M4	Integration, display, then demo preparation	The page shows everything the demo needs to show

Feature freeze in the early afternoon of day three

After the freeze: bug fixes only. No new sources, no new dimensions, no new generation formats.

The last hours are for running real profiles and fixing what they break, not for adding one more thing.

Everyone underestimates how much the last stretch costs. Freezing early is what makes the demo work.

6. Handoff contracts

Four people, three days, and almost all the risk sits in the seams between them.

What each person owes, and to whom

From	To	The contract
M4	Everyone	A running skeleton with stub endpoints by the end of day one, so nobody is blocked waiting for someone else's component.
M1	M3	A list of claim records in the agreed shape. An empty list is acceptable. A list in the wrong shape is not.
M2	M3	A benchmark file containing required terms, per-dimension targets and rate bands for the niche.
M3	M4	A score object, a ranked gap list with blocking items separated out, and generated title and overview text with source spans attached.

Each contract should exist as a fake file from the first two hours. M3 does not wait for M1 to finish parsing; M3 writes a fake claim list in the agreed shape and builds against it. When M1's real output arrives, it drops in.

Example of the shapes

```
claim      { text, skill_ids[], source_type, source_span,
             tier, weight, observed_date, publishable }

benchmark  { niche, required_terms[], title_formula,
             overview_words, portfolio_min, rate_band,
             dimension_targets{} }

result     { readiness, capped, dimensions{}, blocking[],
             gaps[{dimension, current, target, gain,
             effort_hours, priority}], generated{} }
```

7. Checkpoints

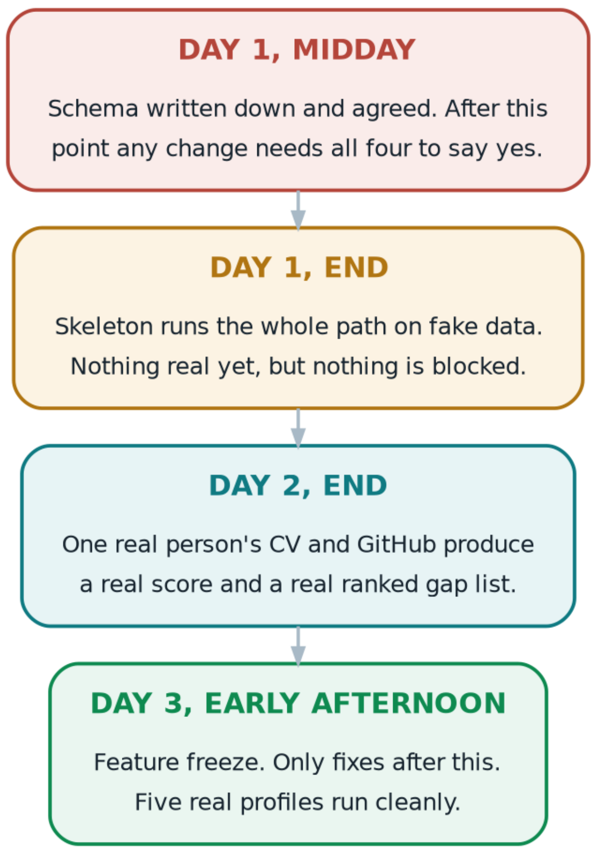

Four moments where the whole team stops and looks at the same screen

Each checkpoint has a decision attached. If the day-two checkpoint fails, the response is to cut scope rather than to work later, because a tired team on day three produces the bugs that break the demo.

If this fails	Cut this
Schema not agreed by midday on day one	Nothing. Stay on it. Everything downstream is worse if this is wrong.

If this fails	Cut this
Skeleton not running by end of day one	Drop the Upwork paste source. Two sources are enough to demonstrate the idea.
No real profile through by end of day two	Drop generation. A score and a ranked gap list still proves the concept.
Integration failing on day three morning	Demo on one prepared profile instead of five, and say so plainly.

8. What will probably go wrong

Risk	Why it happens	What to do
Schema changes on day two	Someone finds a field they need that nobody thought of	Add a field, never rename or remove one. Additions do not break other people.
CV parsing eats a day	Real CVs are far messier than test CVs	Fix the format list on day one. Reject the rest with a clear message.
Benchmark takes longer than planned	Reading thirty profiles carefully is slow work	Cap at thirty. If it is still running late, twenty with honest notes beats thirty rushed.
Everything integrates on the last afternoon	Each person tested only their own piece	The skeleton exists precisely to prevent this. Wire pieces in as they land, not at the end.
Generated text says things the evidence does not support	The span check was left until last	Build the validator with the generator, not after it. It is two hours, and it is the whole point of the product.
Someone finishes early	Scope was uneven	They join integration or run more real profiles. They do not start something new.

9. Definition of done

The sprint succeeded if all of this is true

Five real freelancer profiles have been through the system end to end.

Each produced a readiness score with the cap applied correctly when evidence was thin.

Each produced a ranked gap list with blocking items shown separately from ranked ones.

Each produced a rewritten title and overview, and every number in them can be traced back to a source span on request.

At least one person outside the team has looked at their own report and said something the team did not already know.

The last item is the real test. The first four are engineering; the fifth is whether the product is worth building. A working pipeline that produces reports nobody finds useful is a failed sprint even if every component works.

One line for the team

Two hours on the schema, a running skeleton by the end of day one, and a hard freeze on day three. Those three decisions matter more than anything else in this plan.